

CS774 AI & Ethics

Alice Oh

Aug 31, 2020

Class Objectives

- To understand these issues about AI & Ethics from research perspective
 - Why it's important to consider ethics in developing AI models & systems
 - What research has revealed about ethics-related problems in existing AI models & systems
 - What CS researchers have tried to mitigate these problems
- To replicate & improve on AI & Ethics research
 - Take a paper that takes a computational approach to mitigate an ethics problem
 - Replicate the paper, as much as possible
 - Try something to improve the paper by using a different dataset, a different model, and/or posing a modified research question

Class Tasks

- Read
- Write
- Present
- Discuss
- Research

Class Tasks

- Read
 - We'll have many reading assignments, mostly but not all from CS venues
- Write
 - For some of these papers, I will assign “discussion questions”
 - Write and submit your answers to these questions through Google form/doc
- Present
 - You will present (in a team of 2 students) 2 papers for discussion for one class
- Discuss
 - For most of the classes, you will participate in discussions based on the reading and discussion questions
- Research
 - For the class project, you (in a team of 2 students) will do research based on recent AI & Ethics research papers

Guest Lecturers

- We have 4 guest lecturers
 - Dirk Hovy (Bocconi University)
 - Shakir Mohamed (DeepMind)
 - Joanna Bryson (Herstie School)
 - Kyunghyun Cho (NYU)
- For some, time/date has been fixed, and they are not our normal time
- For these, I will record the lecture and put it on KLMS

First Lecture Survey

- <https://forms.gle/NTZx5k72Q1RgyDMq8>
- You must submit this by the end of Wednesday, Sept 2 (23:59)
- This survey has three questions
 - Describe your programming skills
 - Write which college-level AI/ML class you've taken
 - Accept the rules of academic integrity for this class

Reading for next week

- Big Data's Disparate Impact (Part I only) by Barocas and Selbst
- PDF will be on KLMS